



Implementing OOPs

Tobias Wrigstad
`tobias.wrigstad@it.uu.se`



Features of OOPs

- Polymorphism and subtyping
- Dynamic binding—virtual calls need run-time support
- Run-time type testing
- Inheritance

Only consider class-based OOPs. Prototype-based roughly the same.

Outline

- **Field access**
- Method calls
- Calls through interface types
- Options for untyped languages
- Call-site optimisation techniques

Simplified World View

- Accessing C struct

$x.f$ # address of x + compile-time calculated offset of f in x 's type

- Accessing instance variable

$x.f$ # location (existence) of f depends on run-time type of x 's value

Unified access is desirable

- Simple—can use same method of access everywhere
- Not as simple as records

Inheritance

Multiple inheritance

Separate compilation


Dynamic class loading

Prefixing for unified access

```
class Point2D {  
    int x;  
    int y;  
}
```


```
class Point3D ◁ Point 2D {  
    int z;  
}
```

Point2D p = Point3D()
p.y # can be translated to ~ *(p+1)




Prefixing with multiple inheritance


```
class Point2D {  
    int x;  
    int y;  
}
```




```
class Coloured {  
    colour c;  
}
```



class ColouredPoint \triangleleft Point2D, Coloured {}




which one?



Prefixing with multiple inheritance (cont'd)


Coloured p = ColouredPoint()

p.c # works fine, *(p+0) still denotes a colour c




Point2D p = ColouredPoint()

p.x # breaks! *(p+0) is a colour, not an integer



Point2D p = ColouredPoint()

p.c # breaks! *(p+0) is an integer, not a colour



C++: pointer shifting

- Modify pointer address whenever type of variable changes


Pick any of the possible embeddings

- Keeps field access a constant-time operation (*good*)
- Explicit and implicit casts gets a run-time cost (*bad*)
- Tricky if type information gets lost (e.g., void* pointers) (*bad*)

ColouredPoint cp = ColouredPoint()

Point2D p = cp

Coloured c = cp



C++: pointer shifting

- Modify pointer address whenever type of variable changes

Pick any of the possible embeddings


- Keeps field access a constant-time operation (*good*)
- Explicit and implicit casts gets a run-time cost (*bad*)
- Tricky if type information gets lost (e.g., void* pointers)

Needs RTTI!

```
ColouredPoint cp = ColouredPoint()
```

```
Point2D p = cp
```

```
Coloured c = cp
```



Limitations

- Does not work in an untyped setting
- Subtyping the only way to extend a class layout
 - Does not work with "open classes"
 - Does not work if class layout can be modified dynamically
- For untyped/dynamic/... languages
 - Store fields in a hash table, loads and stores are hash table accesses

(talk more about this later)

Not the whole story

- Check access control at run-time (e.g. due to separate compilation)

Store a table of flags for variables in a class

Perform expensive access control check

- When does offset calculation happen?

Load time—might trigger propagating inclusion

Run-time—perform expensive offset calculation

- Opportunity for optimisation, if classes are invariant

Direct second access after slow first-time check

JIT:ed code can omit checks and use calculated offsets

Object Layout

- Languages with GC, RTTI, etc. will use additional overhead per object, e.g.,
 - Forward pointer space for copying GC, mark bits, etc.
 - Pointer to object's class
 - Sometimes, object can be broken up in slices (e.g., for fragmentation-sensitive applications)
 - Monitor for storing a lock
- Push as much shared information into the class
- Java and C++ are about equally efficient wrt. object layout (except for POD)
- Dynamic languages generally more space demanding

Example: JRuby

- JRuby is considered an efficient implementation of Ruby on the JVM
- Ruby is a dynamic language, fields are ultimately stored in Java hash maps
- Empty JRuby object uses ~72 bytes
 - plus 40 bytes per variable for a 32 bit VM
 - plus 64 bytes per variable for a 64 bit VM
- Compare with an empty Java object that should use <12 bytes

Outline

- Field access
- **Method calls**
- Calls through interface types
- Options for untyped languages
- Call-site optimisation techniques

Simplified World View (cont'd)

- Calling functions and procedures:

foo(y) # location of foo can be determined at compile-time (link-time)

Allows inlining to reduce call-time overhead, etc.

- Calling closures:

foo(y) # push y onto stack, jump to address of foo (~ish)

- Calling methods:

x.foo(y) # which foo depends on run-time type of x's value

Inheritance may require class tree search every call (expensive!)

Subtype polymorphism and late binding/class loading makes efficiency difficult to achieve

Method call in untyped OOPs

- Search inheritance hierarchy from x's class for foo/1

Very slow lookup time! (function of #classes and #methods)

- Use a hash table in each object `x.foo(y) ~> push x,y + jmp x.get(foo/1)`

Still much slower than a procedure call!

- Make an entry for each method in the object just like a field

Fast, constant-time dispatch (load + jump)


Very large objects

- Optimisation: share method entries for objects of same class in *vtables*

Much smaller object for the cost of one extra indirection

Vtables for efficient dispatch

- Virtual tables
- Complication due to multiple inheritance




p.getY()

```
void* header = p-1;  
void* vtable = *header;  
int (*getY)() = vtable+1;  
int temp = getY(p); // this
```

Vtable prefixing with single inheritance

```
class Point2D {  
    int x, y;  
    int getX() ...  
    int getY() ...  
}
```

```
class Point3D ◁ Point2D {  
    int z;  
    int getZ() ...  
}
```




Vtable prefixing with multiple inheritance

```
class Point2D {  
  int x, y;  
  int getX() ...  
  int getY() ...  
}
```

```
class Coloured {  
  colour c;  
  colour getC() ...  
}
```

```
class ColouredPoint < Point2D, Coloured { }
```

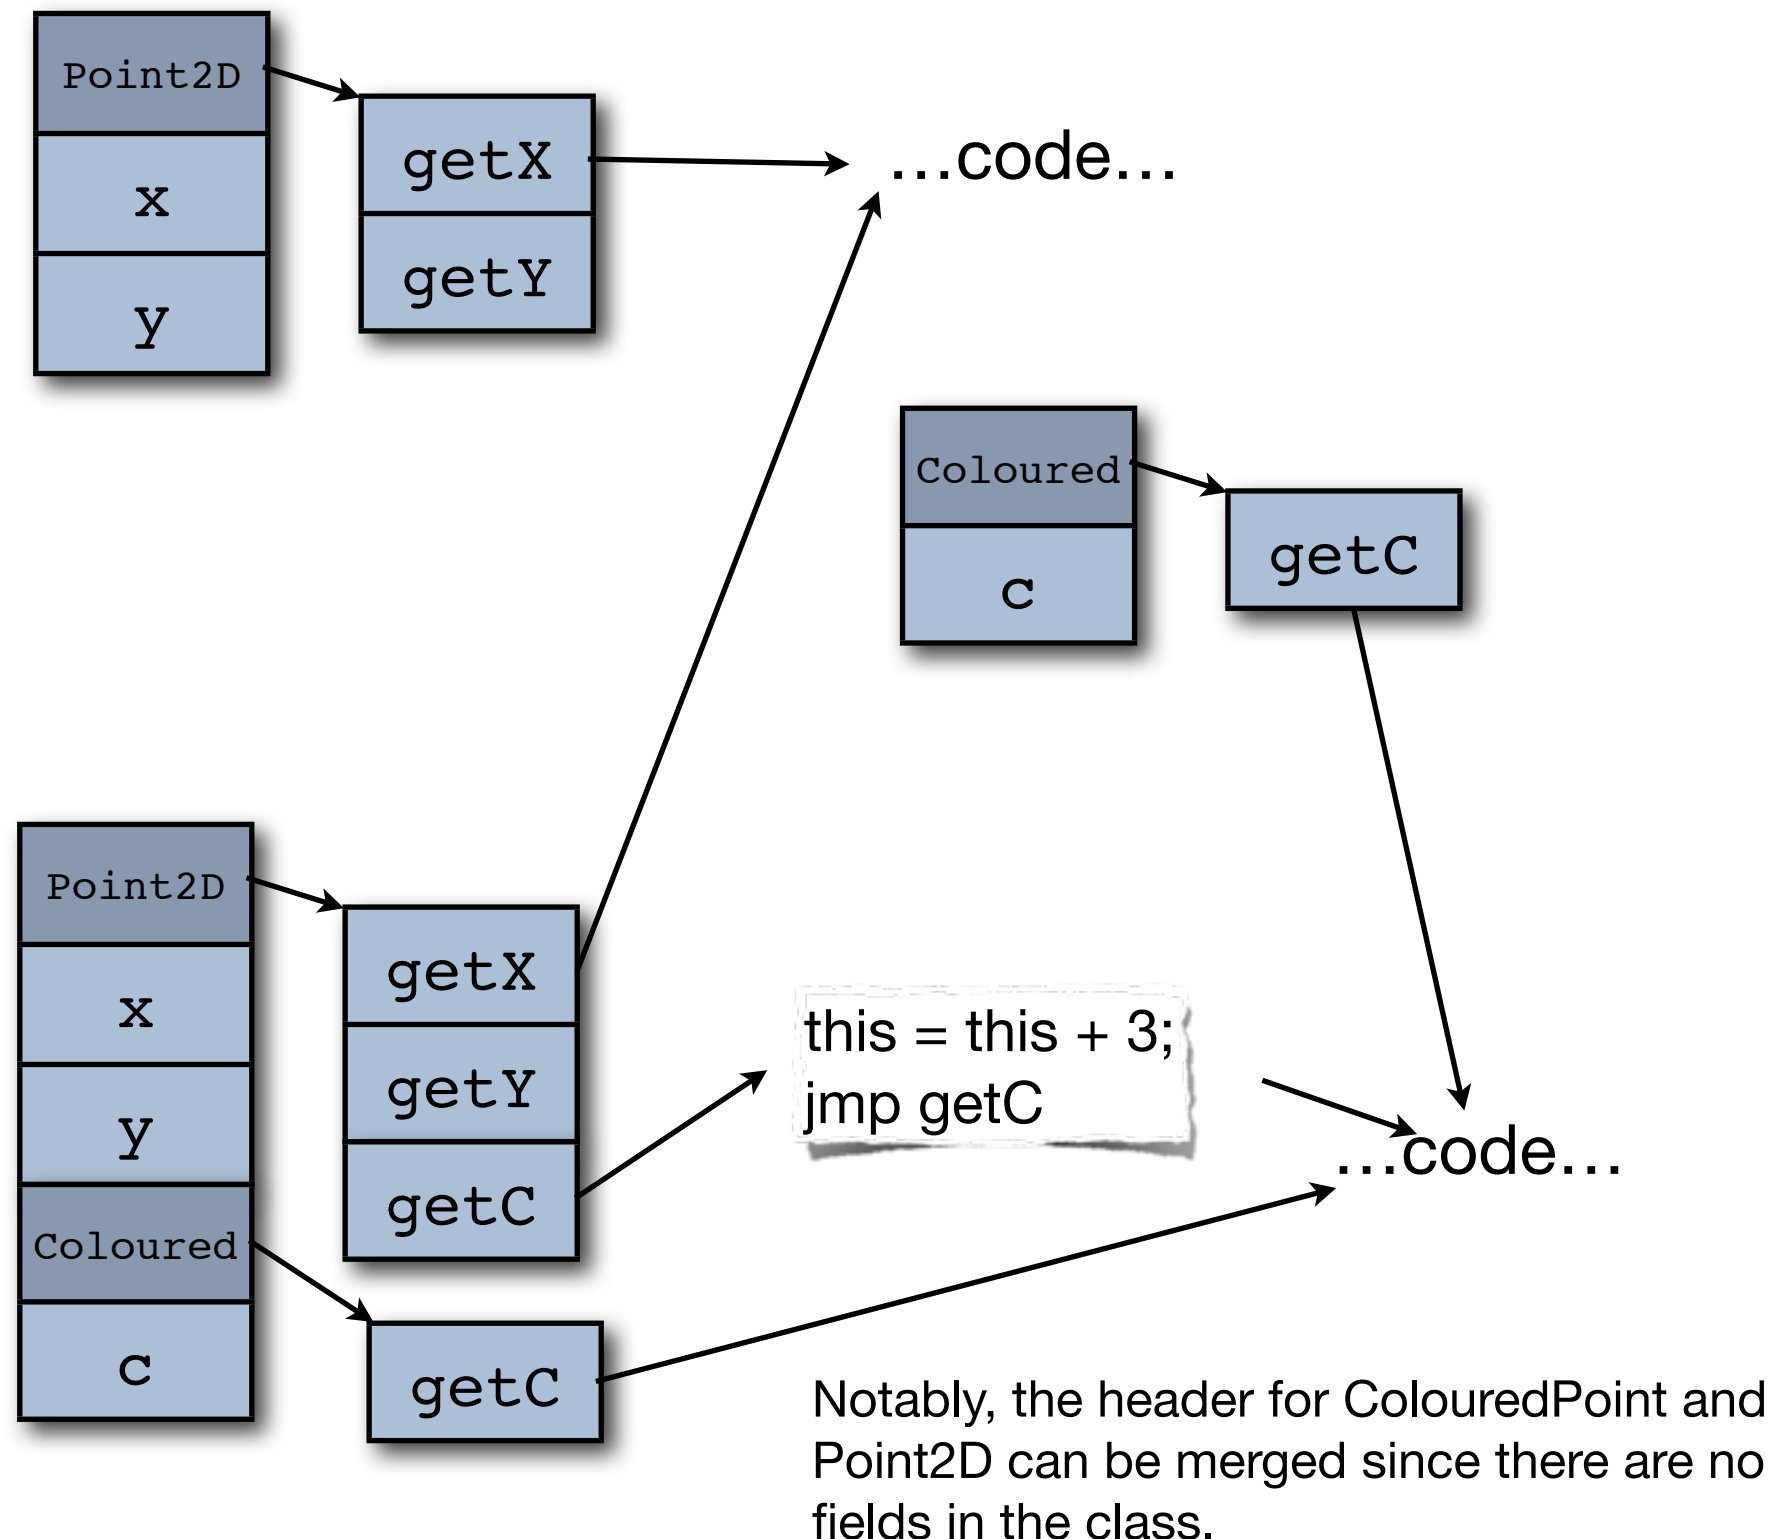


Vtable prefixing with multiple inheritance

```
class Point2D {  
  int x, y;  
  int getX() ...  
  int getY() ...  
}
```

```
class Coloured {  
  colour c;  
  colour getC() ...  
}
```

```
class ColouredPoint < Point2D, Coloured { }
```



Outline

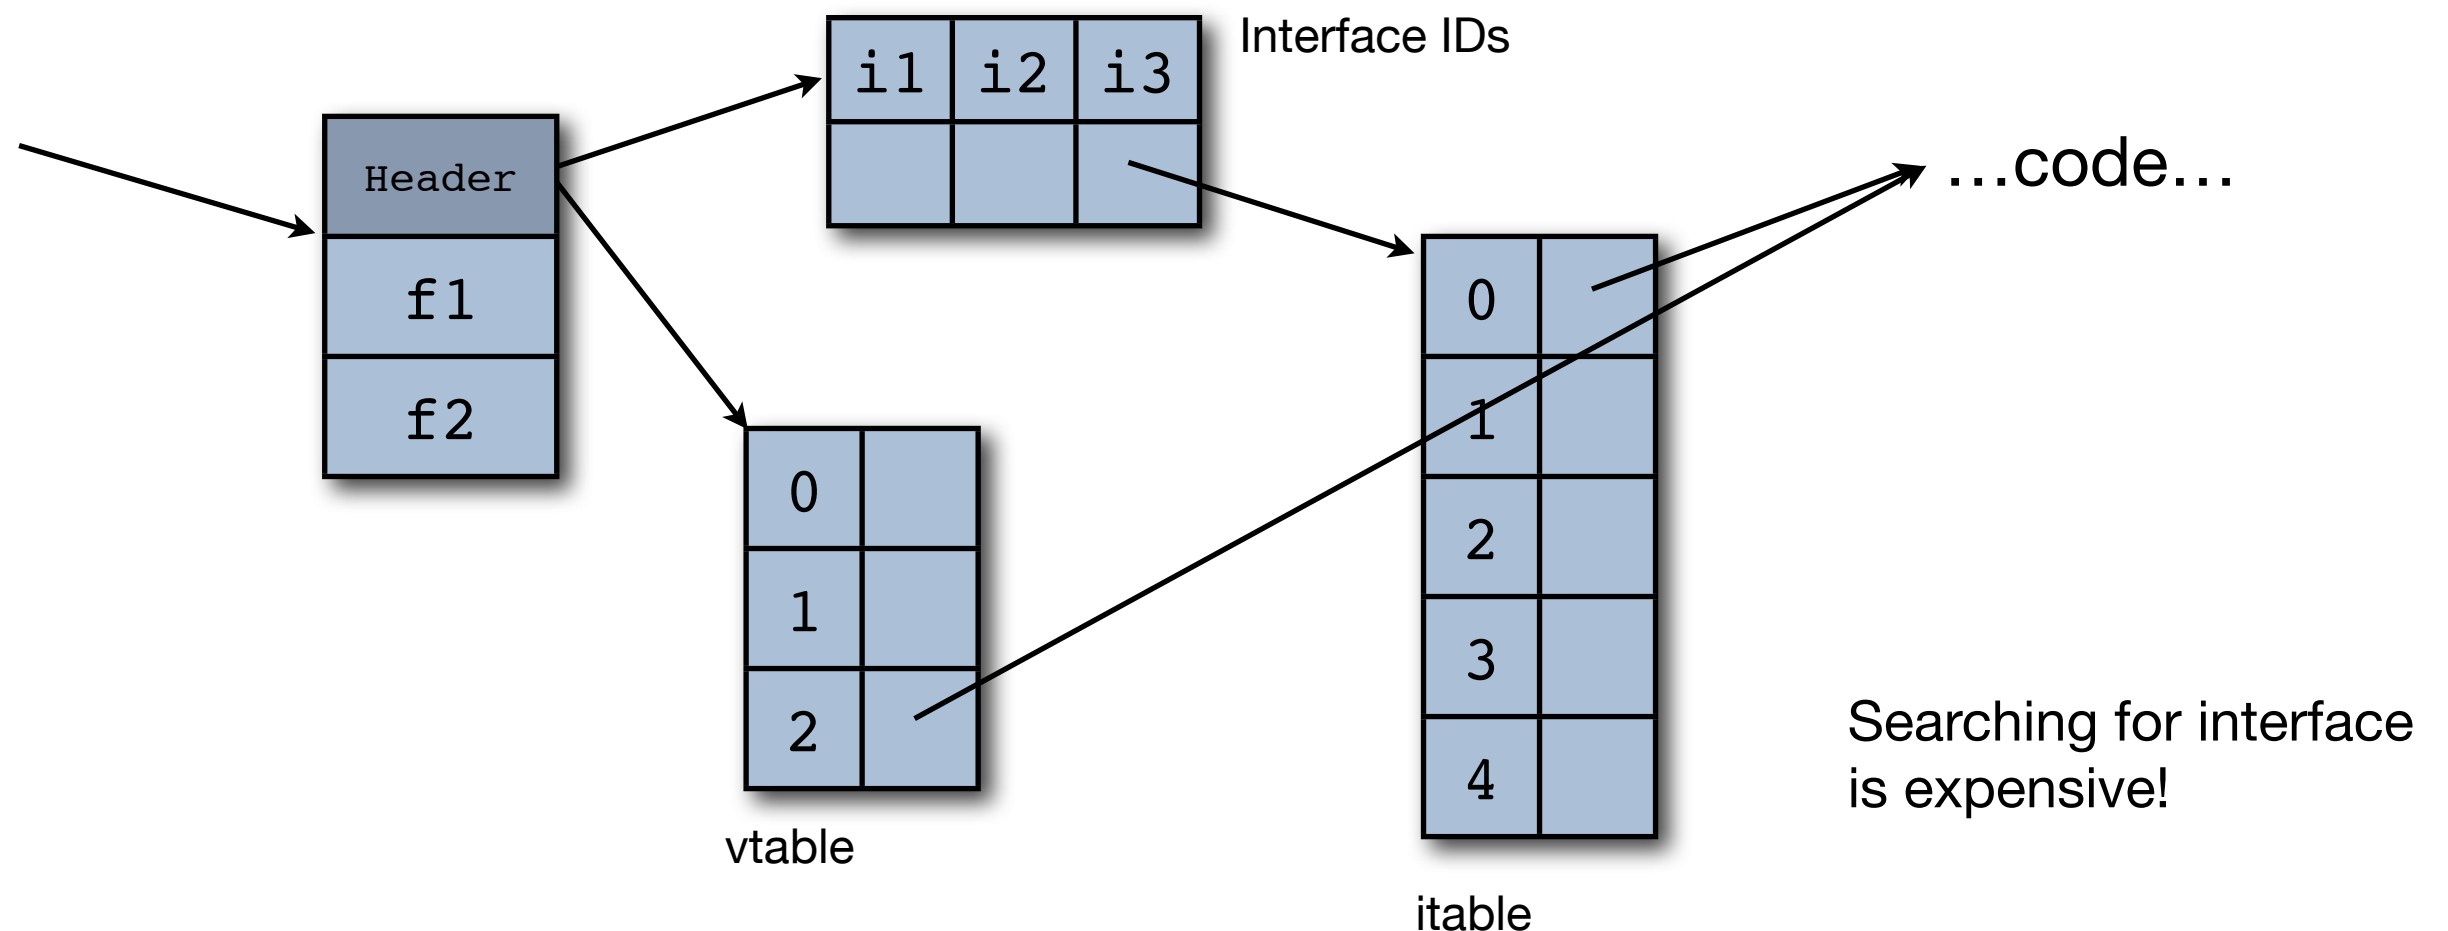
- Field access
- Method calls
- **Calls through interface types**
- Options for untyped languages
- Call-site optimisation techniques

Invocation through interface types

- Observation: Impossible to achieve uniform access through interface types
- **Technique 1:** translate interface offsets to implementing class offsets

Each class has a dispatch "itable" for each interface it implements

On call: search for correct itable by some interface id and use it for dispatch

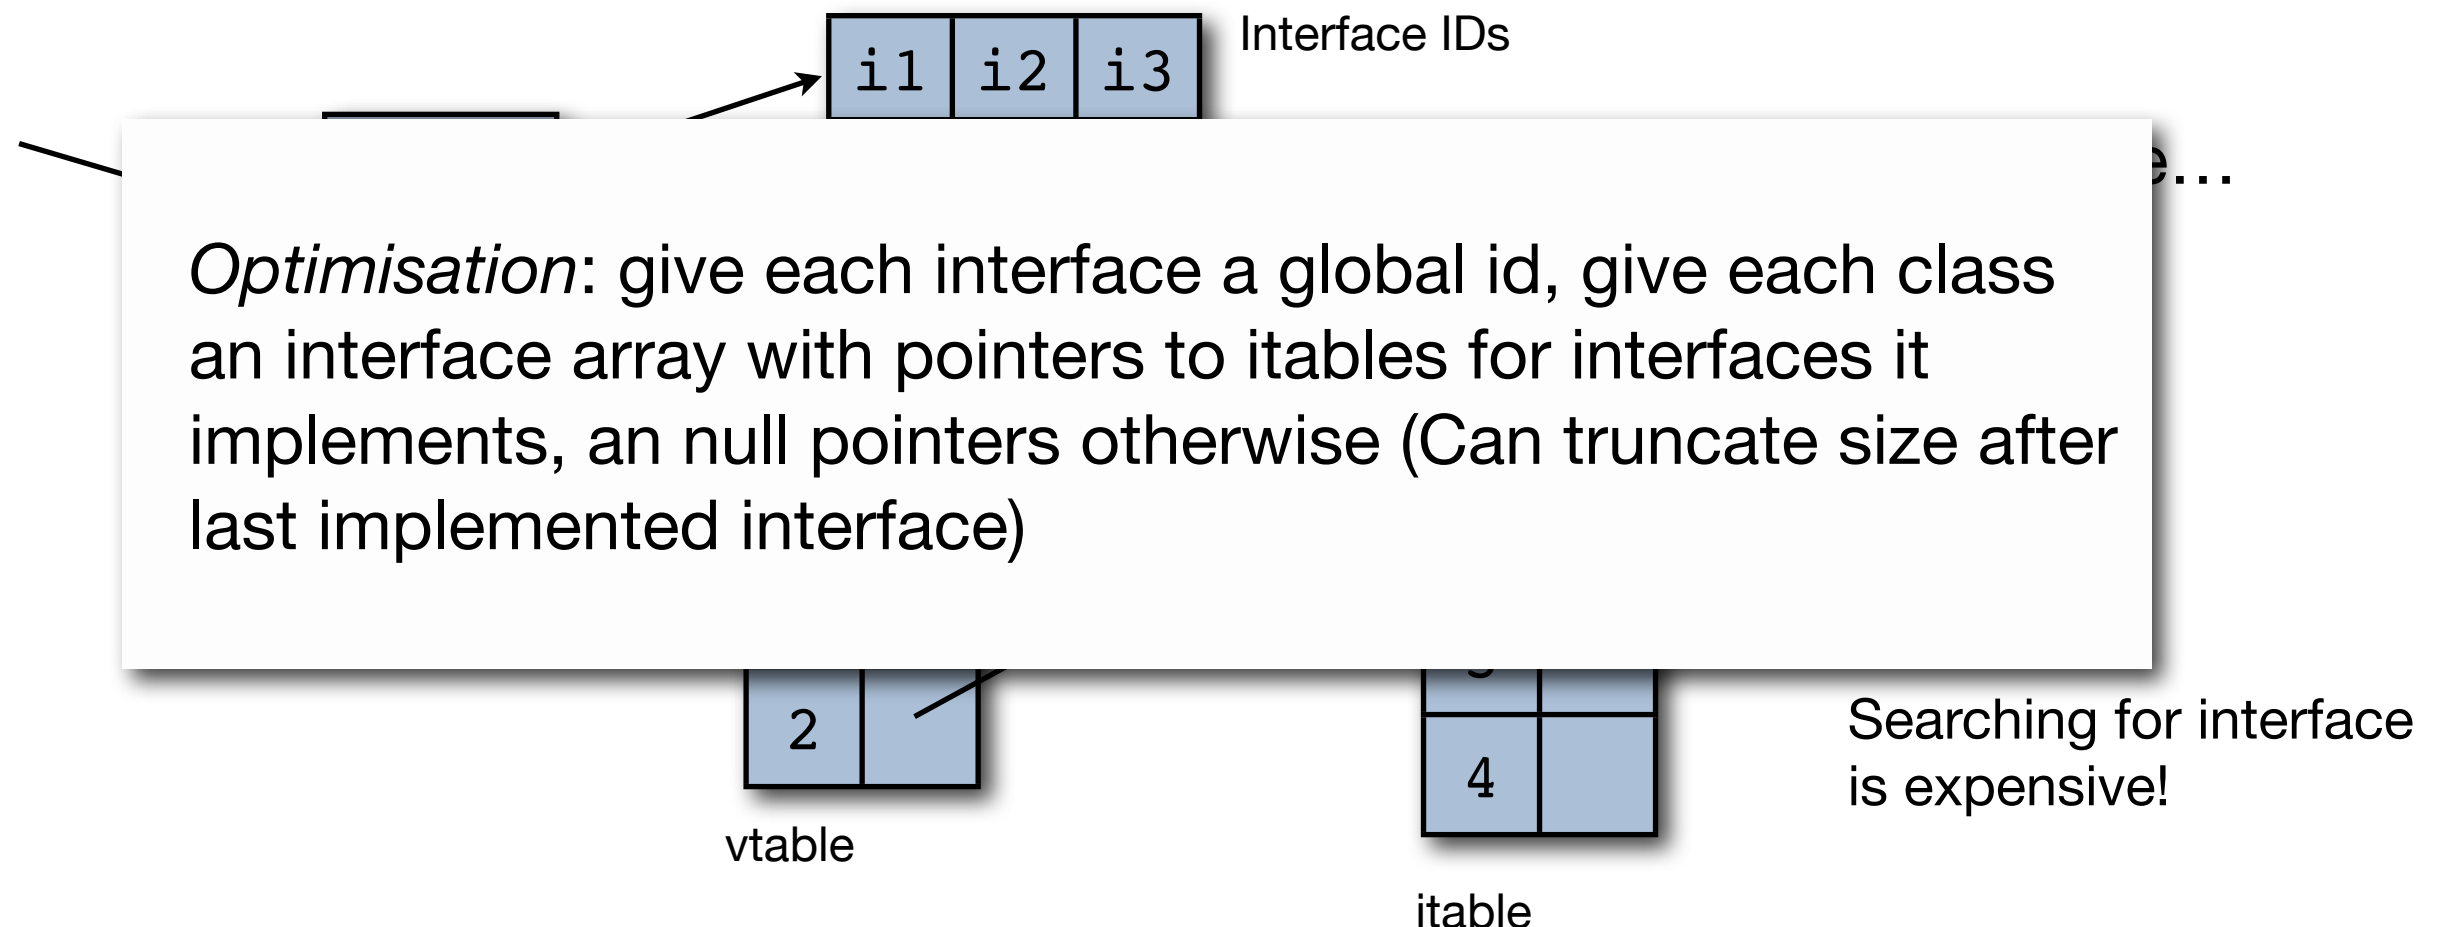


Invocation through interface types

- Observation: Impossible to achieve uniform access through interface types
- **Technique 1:** translate interface offsets to implementing class offsets

Each class has a dispatch "itable" for each interface it implements

On call: search for correct itable by some interface id and use it for dispatch



Invocation through interface types (cont'd)

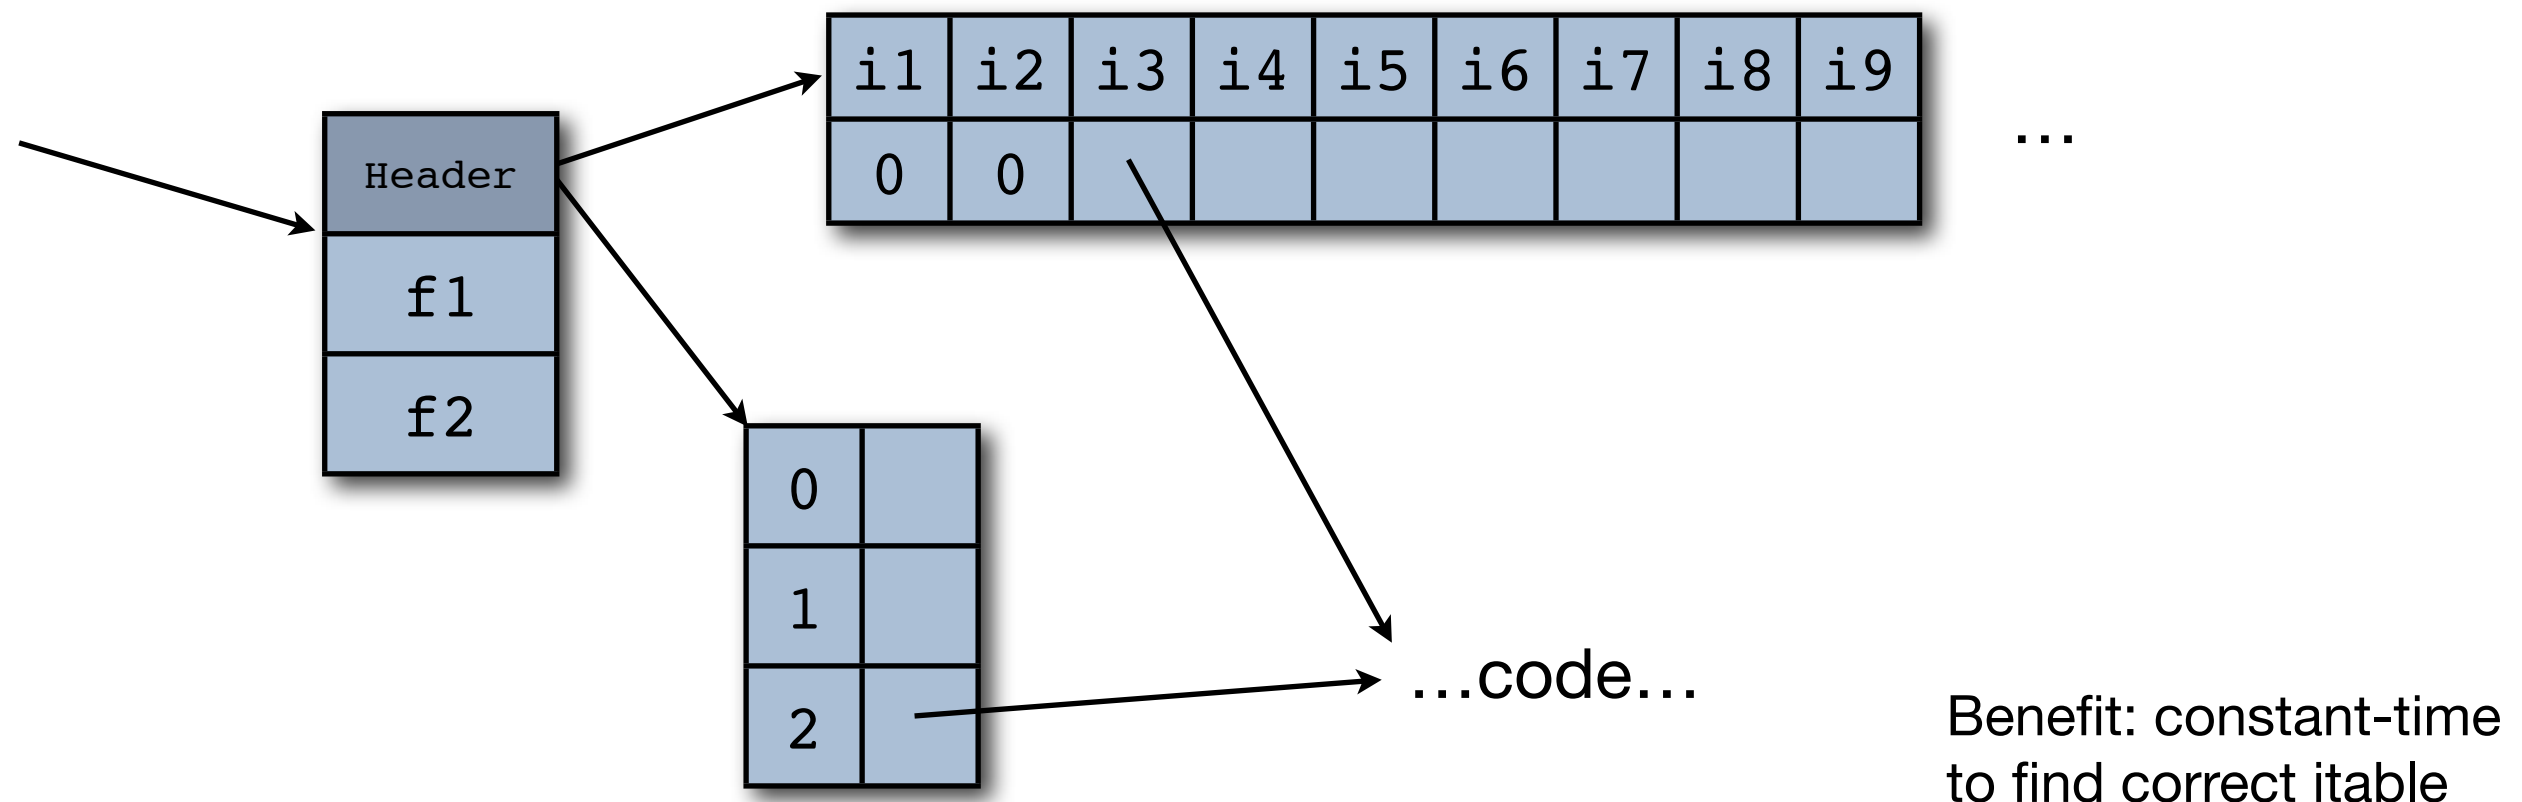
- **Technique 2:** "Selector-indexed tables"

Give each interface method a numerical id (e.g., at load-time)

Give each class an itable for *all* methods in *all* interfaces

Dispatch becomes additional indirection—lookup in the selector index table

Fast but very costly wrt. space



Invocation through interface types (cont'd)

- **Technique 2:** "Selector-indexed tables"

Give each interface method a numerical id (e.g., at load-time)

Give each class an itable for *all* methods in *all* interfaces

Dispatch becomes additional indirection—lookup in the selector index table

Fast but very costly wrt. space

Optimisation: Use graph colouring to re-use same numerical index for methods that may never be called on the same classes.

Can greatly reduce size of all itables.

1	
2	

...code...

Benefit: constant-time
to find correct itable

Invocation through interface types (cont'd)


- **Technique 2:** "Selector-indexed tables"

Give each interface method a numerical id (e.g., at load-time)

Give each class an itable for *all* methods in *all* interfaces

Dispatch becomes additional indirection—lookup in the selector index table

Fast but very costly wrt. space



i ₁	i ₂	i ₃	i ₄	i ₅	i ₆	i ₇	i ₈	i ₉	i ₁₀
----------------	----------------	----------------	----------------	----------------	----------------	----------------	----------------	----------------	-----------------

Observation: We can get more space-efficient by using a small selector table. But hashing might produce duplicates—same index for different methods.

Solution: store pointer to code that switches on receiver class type to find the correct implementation.

Class type can be a *hidden parameter* [Alpern et. al., OOPSLA 2001] constant-time
to find correct itable

Outline

- Field access
- Method calls
- Calls through interface types
- **Options for untyped languages**
- Call-site optimisation techniques

Alternatives for untyped code

- No type information to base field, vtable (or itable) offsets from
- Storing closures in hash tables is not space efficient for *methods* (but for fields)
- Naïve implementation:
 - Search receiver's class' hierarchy for method signature, and call
 - Slow, terrible worst-case times
- Possible to use static table internally, esp. with JIT compilation

Global dynamic table

- Global cache in the form of a hash table indexed by class + signature

Translate invocations into lookups in hash table

Cache miss: perform expensive search through class hierarchy, then update cache

- Flush (part of) the cache as a result of reflective operations that change classes
- Space costs are reasonable
- Overhead is reasonable and table is constructed incrementally

Global dynamic table

- Global cache in the form of a hash table indexed by class + signature

Translate invocations into lookups in hash table

Cache miss: perform expensive search through class hierarchy, then update cache

- Flush (part of) the cache as a result of reflective operations that change classes
- Space costs are reasonable
- Overhead is reasonable and table is constructed incrementally

Is it effective?

OK average call time, bad worst-case call time

One dispatch table per message names

- Create a separate table per unique signature mapping classes to methods
- Each call site can statically know what table to consult
- Performance is better than the single global table

Especially if methods names are relatively unique

(Smalltalk names fare quite well here)

- Per-signature dispatch tables can be constructed incrementally

Outline

- Field access
- Method calls
- Calls through interface types
- Options for untyped languages
- **Call-site optimisation techniques**

Inlining

- Difficult to do in flexible programs

Analysis of e.g., possible run-time binding is limited by dynamic loading

- Java

Can possibly inline **final** and **static** methods

JITing allows more aggressive inlining

```
class A {  
    int foo(int x, int y) { return x+y; } // can be inlined by HotSpot if  $\nexists B <: A$   
}
```

- Dynamic class loading requires remembering JITed methods and storing their prerequisites (e.g. $\nexists B <: A$ above)

Check prerequisites and possibly "retire" (unoptimise) compiled code on class loading

Call-site optimising through caching

- Useful esp. for untyped code and interface calls
- Techniques addressed here:

Inline Caching [Deutsch and Shiffman, 1984]

Polymorphic Inline Caching [Hölzle et al., 1991]

Inline Caching [Deutsch and Shiffman 1984]

- Each call site has a single-element lookup cache

Remember what actual method was called for class of last receiver

Next call, if same receiver we can get method immediately from cache

Cache miss: slow-path through lookup, update caches

- Efficient implementation through self-modifying code

x.m(...)

```
c = x.class  
am = c.search(m/1)  
jump am
```



next call

```
switch (x.class) {  
  case c: jump am; break  
  default:  
    c = x.class  
    am = c.search(m/1)  
    jump am }  

```

Inline Caching [Deutsch and Shiffman 1984]

- Each *Is it effective?*
 - Re *Smalltalk: 90-95% cache hit frequency and ≈ 4 instructions for fast path.*
 - Ne *Slow path real slow though.*
 - Ca *Polymorphic and megamorphic call sites terrible performance!*
- Efficient implementation through self-modifying code

x.m(...)

```
c = x.class  
am = c.search(m/1)  
jump am
```



next call

```
switch (x.class) {  
  case c: jump am; break  
  default:  
    c = x.class  
    am = c.search(m/1)  
    jump am }  
}
```

Polymorphic inline caching [Hölzle and Ungar, 1991]

- Handles polymorphic and megamorphic call sites

Extension is simple: use a multi-element cache

Allows relatively fast dispatch for polymorphic call sites

If several classes are equally common, performance degrades

Can get large space overhead (esp. for megamorphic call sites)

call when x is a Bar

```
switch (x.class) {  
  case Foo: jump m1 break;  
  default: ... # lookup + install  
}
```



```
switch (x.class) {  
  case Foo: jump m1 break;  
  case Bar: jump m2 break;  
  default: ... # lookup + install  
}
```

Polymorphic inline caching [Hölzle and Ungar, 1991]

- Handles polymorphic and megamorphic call sites

Extl *Extensions: change case ordering based on hit frequency.*

Allk *But will it earn back the incurred run-time overhead?*

If s

Can get large space overhead (esp. for megamorphic call sites)

call when x is a Bar

```
switch (x.class) {  
  case Foo: jump m1 break;  
  default: ... # lookup + install  
}
```



```
switch (x.class) {  
  case Foo: jump m1 break;  
  case Bar: jump m2 break;  
  default: ... # lookup + install  
}
```

Polymorphic inline caching [Hölzle and Ungar, 1991]

- Handles polymorphic and megamorphic call sites

Ext *Extensions: change case ordering based on hit frequency.*

All *But will it earn back the incurred run-time overhead?*

If s *... # lookup + install*

Can get large space overhead (esp. for megamorphic call sites)

Improvement: use binary search instead of linear switch.

sw *Requires global knowledge (to map classes to integer ids)*

ca *Little additional overhead*

de
}

default: ... # lookup + install
}

JRuby use of IC [Bini et al.]

```
public IRubyObject call(IRubyObject caller, IRubyObject self, IRubyObject arg1) {  
    RubyClass selfType = pollAndGetClass(self);  
    if (CacheEntry.typeOk(localCache, selfType)) {  
        return localCache.method.call(self, selfType, methodName, arg1);  
    }  
    return cacheAndCall(caller, selfType, self, arg1);  
}
```

(Simplified to fit on screen)

(Notably not polymorphic)

References

1. Craig Chambers, *Efficient Implementation of Object-Oriented Programming Languages*, Tutorial, OOPSLA 2000
2. Bowen Alpern, Anthony Cocchi, Stephen J. Fink, David Grove, and Derek Lieber, *Efficient Implementation of Java Interfaces: Invokeinterface Considered Harmless*, OOPSLA 2001
3. Urs Hölzle and Ole Agesen, *Dynamic vs. Static Optimization Techniques for Object-Oriented Languages*, in Theory and Practice of Object Systems 1(3), 1995
4. Urs Hölzle, Craig Chambers and David Ungar, *Optimizing Dynamically-Typed Object-Oriented Languages With Polymorphic Inline Caches*, ECOOP 1991
5. L. Peter Deutsch and Allan M. Schiffman, *Efficient implementation of the smalltalk-80 system*, POPL 1984
6. Stefan Matthias Aust et al., *JRuby 1.4.0. source code*, retrieved Feb 2010